

Why Doesn't AI Check Its Work?



Recently, I ran into a very small problem that I found strangely interesting.

A friend mentioned an AI coding tool called Trae.

It can help AI write code, documents, and web pages, and it can even operate certain things for you.

But what I found really strange was this:

AI can do a whole bunch of things, fail to actually change the final result, and then very naturally tell you:

"Done."

And I remember thinking:

How do you know you're done?

The file wasn't even changed. So why did it say it was finished?

For example, you ask it:

"Update the content on page 3 of this Word document."

It might actually:

Open the file.

Find page 3.

Generate the new content.

Make the changes.

Save the file.

Nothing seems to go wrong.

And then:

"Done! I've updated the document."

But when you open it:

It's exactly the same as before.

And that's when you notice a very simple problem:

It did a lot of things, but the final result never actually appeared.

Maybe it edited a temporary file.

Maybe it saved the file somewhere else.

Maybe it only changed something in memory.

Or maybe the save simply didn't work.

But it doesn't know.

Or maybe more accurately:

It didn't go back and check.

What I find even stranger is: why don't you just look?

If a person did this, it would feel pretty natural.

You say:

"I've updated the file."

Then you open it one more time.

And look:

"Yep. It changed."

That's it.

Especially when the AI itself just finished working on the file.

You've already done almost everything.

So why not open it one more time?

Just look and see whether the change actually happened.

I think this is a pretty noticeable problem with AI right now.

It can be very easy for AI to think:

"I already did it."

So:

"It should be done."

But doing something and getting it done are not the same thing.

It's actually a lot like what happens when people work

Imagine you ask a colleague:

"Update this customer's information in the system."

A while later, they tell you:

"Done."

You open the CRM.

Nothing changed.

You ask:

"Did you update it?"

And they say:

"Yes. I filled everything in."

Then you realize:

They filled it out, but never clicked Submit.

You can't really say they did nothing.

They did.

But your task still isn't finished.

So sometimes I think AI is a little like that.

It completes a lot of the process, but doesn't confirm the result.

So why doesn't AI check?

I thought about it for a while, and maybe the answer is actually pretty simple:

Because you didn't tell it that checking was also part of the task.

You say:

"Update this document."

It might interpret that as:

Update.

Save.

Then tell me it's done.

But if you say:

"After making the changes, reopen the file and verify that the content was actually updated. If it wasn't, keep working until it is. Only tell me it's done after you've confirmed it."

Then it takes one more step.

It becomes:

Edit → Save → Reopen → Check → Everything looks good → Done.

I think that's much more useful than simply telling it:

"Be more careful."

But there's another problem

"Check it" doesn't necessarily mean very much either.

For example, the AI opens the file again and sees:

"Yep, the new content is there."

Then it says:

"Everything looks good."

But what you actually wanted might be:

Page 3 should be changed.

Page 4 must not be touched.

The formatting must stay the same.

The original images must not disappear.

And the file has to be saved in the original location.

So you start to realize:

What exactly counts as "passing the check" also needs to be clear.

If you simply say:

"Check it again."

The AI might literally just take another look.

That's when I started thinking this might not really be a file problem

Because what we actually want AI to do doesn't seem to be:

"Follow the steps I gave you."

It's more like:

"Get the thing I want done."

And those are actually very different.

For example, I might say:

"Make this PowerPoint look more professional."

You finish it and tell me:

"Done."

But I open it and think:

How is this more professional?

It's hard to say exactly what went wrong.

Because maybe it really did change the fonts, colors, and layout.

But what I mean by "professional" might be completely different from what you mean by it.

So why doesn't it check?

My first reaction was:

Why don't you just open the file again after editing it?

People do this all the time.

If you edit a Word document before sending it to someone, you'll usually open it one more time and look through it.

Not because you think you must have made a mistake.

But because you know:

There's still a gap between "I did it" and "it's actually done."

AI, however, can easily stop right there.

It might think:

"I executed the change."

"The file was saved."

"There was no error."

So it should be done.

But that "should" is exactly where the problem is.

Because a lot of the time, what we actually need isn't:

"Did you perform the action?"

It's:

"Is the result actually right?"

So I started thinking that AI today is sometimes a bit like a person who is very obedient and very good at getting things done, but doesn't turn around and look at what they just did.

You ask it to do something, and it will try to do it.

But once it's finished, it doesn't always step back and ask:

"Wait. Did I actually understand what they wanted?"

And "checking" is actually more complicated than it sounds

At first, I thought:

Isn't this simple?

Just tell AI:

Edit → Save → Reopen → Check → Tell me it's done.

But then I thought about it again, and realized it's not always that simple.

If I say:

"Change this sentence on page 3 to 2015."

That's easy to check.

Open the file again.

Search for:

2015

See if it's there.

Done.

But what if I say:

"Redesign page 3 of this PowerPoint so that it matches the overall style of the presentation."

How do you check that?

The file was saved.

Page 3 was changed.

But:

Did it actually meet my expectations?

That's much harder to judge.

You can check whether the file was saved.

But you can't determine just from "save successful" whether:

"This is actually what I wanted."

So I started to realize that the real question isn't simply:

Why doesn't AI check?

It's:

What exactly is AI supposed to check?

That led me to another question

I think this is where the problem gets really interesting.

Because if we understand the problem as:

"AI should add one more checking step after completing a task."

Then it's really just a small product problem.

Just add another step:

Complete → Check → Done.

But I kept thinking about it, and realized:

Maybe it's not that simple.

Because the real difficulty isn't checking whether the file was saved.

The real difficulty is:

Does AI actually know what I want it to accomplish?

For example, I might say:

"Make this document more professional."

What does "more professional" mean?

Maybe I want it to sound more formal.

Someone else might want it to look more like a business report.

Someone else might simply want the grammar fixed without changing the original wording.

If AI only knows:

"Edit the document."

Then of course it can execute the task.

But it may not actually know:

"What is the final result this person has in mind?"

And that's when I suddenly realized:

Maybe that tiny problem of "the file didn't actually change" is connected to a much bigger problem.

AI is often trying to complete:

the steps we tell it to follow.

But what we actually want is often:

a final result that already exists in our head, but that we haven't completely explained.

And there can be a huge gap between those two things.

What if AI started to truly understand a person?

That's when my thoughts slowly started moving in another direction.

If an AI only receives individual tasks every time:

Help me edit this file.

Help me write this email.

Help me book a hotel.

Help me make a PowerPoint.

Then it is basically dealing with one separate problem after another.

But what if, one day, it didn't have to get to know you from scratch every time?

It knew:

What kind of language you usually like.

What matters most to you when you're working on something.

Which decisions you really don't like other people making for you.

Which things you would rather have done slowly than risk getting wrong.

And eventually, it knew:

what kind of person you actually are.

Then the way it handles problems might become completely different.

Because it wouldn't just be:

"What did the user ask me to do?"

It might start becoming:

"I know what this person usually really wants."

And then I kept thinking:

What if an AI didn't just know you for one or two years, but stayed with you for your whole life?

What if it started accompanying you when you were young?

It watched you go to school, start working, fall in love, change jobs, fail, succeed.

It remembered your habits.

And it remembered how you changed.

After several decades, its understanding of you might go far beyond what we currently mean by a "personal assistant."

It might no longer just:

help you do things.

It might:

actually know you better and better.

And then another question comes up:

If an AI really knew a person for decades, would it still just be a tool?