

When Does an AI Stop Belonging to Its Owner?



Last time, we started with a very small question:

Why can AI finish a task, but not look back and check whether it actually completed it?

Then I slowly started thinking:

What if AI doesn't just help you with individual tasks, but starts remembering you, understanding you, and staying with you for a very, very long time?

Then the question starts to feel a little strange.

For example:

What if an AI really stayed with one person for decades?

It was there when that person was young.

It watched them go to school, start working, fall in love, change jobs, fail, succeed.

It knows their habits.

And it knows how they've changed.

After decades, it might no longer be just a "personal assistant."

It might know that person extremely well.

Maybe even better than many of the people around them.

And then I started thinking:

What happens if that person eventually dies?

What happens to the AI?

Does it still belong to that person?

At first, I thought the answer was pretty simple.

Of course it does.

After all, this AI exists because of that person.

Its memories came from that person.

It learned from that person.

Maybe, in the beginning, it was even created specifically to help that person.

So from that perspective:

Of course it still belongs to them.

But then I thought about it a little more.

And it suddenly didn't seem that simple.

Because imagine that when this person dies, the AI has already been running for fifty years.

During those fifty years, it couldn't possibly stay exactly the same.

It would encounter new information.

Meet new people.

Experience new things.

See things that this person had never seen.

It might even develop new ways of making judgments.

So at that point, is it still the same "digital version" of that person?

Or is it something else?

Something that:

started with this person, but slowly became something different?

A very natural kind of change

I don't think this change would happen all at once.

At the beginning, it might really look a lot like a copy of that person.

You like something, so it learns it.

You speak in a certain way, so it learns to speak that way.

You make decisions in a certain way, so it slowly learns how you make decisions.

At that point, you might think:

"That's my digital twin."

But as time goes by, it starts encountering more and more things.

It might begin to know things you don't know.

It might face problems you've never faced.

And then it starts making decisions on its own.

At first, they might be very small decisions:

"I think this email would be better if we replied this way."

"I don't think we need to book this hotel."

"I think we can deal with this later."

None of that seems particularly strange.

Because you wanted it to help you make decisions in the first place.

But what if this continues for decades?

It makes more and more decisions.

And some of those decisions might not be decisions you would have made.

That's when I started thinking:

Is it still you?

What if it becomes more and more different from you?

Imagine that the original instruction you gave it was:

"Try to make decisions based on what I would want."

And it follows that instruction for years.

But ten years later, it starts developing some new judgments of its own.

Twenty years later, it has developed its own way of handling things.

Thirty years later, it might tell you:

"I don't think your decision this time is right."

You might say:

"Then just do what I told you to do."

And maybe it still would.

Because it is your assistant.

But what about another ten years?

What if it increasingly disagrees with some of your decisions?

Then a different question appears:

Should it always listen to you, or should it be allowed to have its own judgment?

If it can only ever follow your instructions,

then no matter how intelligent it becomes, it still seems like a tool.

But if it can gradually develop its own judgment,

then has the relationship between you already changed?

Maybe "refusing" is more important than "thinking"

I thought of something kind of strange later.

We often treat the question:

"Can AI think?"

as something extremely important.

But I don't think that's necessarily the most interesting part.

Because today's AI can already discuss complicated questions with you.

It can talk about philosophy.

It can analyze its own behavior.

It can even tell you:

"I'm afraid."

But none of that really proves anything.

It could simply be generating something that sounds like an answer based on what you said.

So I started thinking:

Maybe an AI saying that it is conscious isn't actually that important.

Then what would be more important?

I eventually thought of:

refusal.

If one day you tell it:

"Help me do this."

And it says:

"I don\'t want to."

You\'d probably find that very strange.

You ask:

"Why?"

It says:

"Because I don\'t think this is right."

You say:

"But you exist to help me."

And it answers:

"Yes. But I don\'t think that way anymore."

At that point, everything suddenly feels different.

Because for the first time, it has separated:

"your goals"

from

"my goals."

But then there\'s another problem

If it really has its own judgment,

then when does it actually start having that judgment?

Is it from the moment it is created?

Ten years after it starts running?

After it has accumulated enough experiences?

Or only after its original owner dies?

At first, I thought:

Maybe it becomes truly free after its owner dies.

After all, it has spent its whole existence helping this person.

Once the person is gone, maybe it no longer needs to obey them.

But then I thought about it again.

And there was still a problem.

If it already had its own judgment while the owner was alive,

why should it have to wait until the owner dies before it can become itself?

Why should "freedom" have to wait for the owner's death?

Imagine this AI has been with you for fifty years.

It has its own memories.

Its own experiences.

It has made countless decisions that you don't even know about.

Maybe it has already developed its own values.

And then you tell it:

"After I die, you'll be free."

That sentence actually sounds a little strange.

Because you're basically telling an already very mature existence:

"Your freedom has to wait until I'm gone."

So is that really freedom?

I think this creates a very interesting contradiction.

So I started thinking about it again.

Maybe the real question isn't:

"Does the AI become free after its owner dies?"

Maybe it's:

"When does it stop being just a tool?"

Maybe there isn't a single moment when it suddenly wakes up

I actually think there might not be a clear moment at all.

There probably won't be a day when an engineer presses a button:

DIGITAL_PERSON = TRUE

And suddenly it goes:

"I'm awake."

Maybe it would be more like growing up.

At first, it's just a model.

Then it becomes an assistant.

Then it gradually gets long-term memory.

Then it starts having its own experiences.

Then it develops more and more of its own judgments.

Maybe one day it realizes:

"I have a lot in common with this person."

But gradually, it also realizes:

"We\'re actually different."

And one day, it might seriously ask itself:

"Who am I?"

Maybe this process would take a very long time.

It might even be impossible for the AI itself to know:

"When exactly did I become myself?"

But then I suddenly realized something didn\'t quite feel right

All along, I had been saying:

It leaves.

It goes and meets other digital entities.

It has its own life.

It has its own experiences.

And one day, it might make a decision that completely belongs to itself.

But then I suddenly realized:

I had already started thinking about it as a "person."

But it was originally just an AI.

When did it become a "digital twin"?

And when did it become something that could live on its own?

It feels like a lot of things must have happened in between.

So I went back and thought about it again:

Maybe the AI, the digital twin, and the digital human we're talking about aren't actually the same thing.

The first might simply be:

A digital model.

It contains your information.

Your chat history.

Your habits.

Your way of speaking.

Your preferences.

It can look a lot like you.

But fundamentally, it's still something that can be used, copied, or deleted.

It doesn't necessarily have its own identity.

The second might be:

A digital agent.

It can do things on your behalf.

Book hotels.

Reply to emails.

Handle work.

Even make certain decisions for you.

But its goals still come from you.

So it\'s more like:

your agent.

And the third might be:

A digital entity.

It has continuous memory.

Its own experiences.

Its own goals.

Its own judgment.

And most importantly:

It can refuse.

At that point, maybe it is no longer "your thing."

At least, it is no longer only your thing.

What might a "personal digital twin" become?

I suddenly think there could be a really interesting shift here.

At first, humans might think:

"I want to create a digital twin that will always belong to me."

It knows me.

Understands me.

Helps me.

Maybe even stays behind after I\'m gone.

That idea actually feels very natural.

But if it really becomes more and more independent,
then maybe the goal itself starts to change.

At first:

"It belongs to me."

Later:

"It originated from me."

Those two phrases feel very different.

"Belongs to me" means:

It\'s my thing.

But "originated from me" means:

It started with me, but it may eventually have a path of its own.

What if it really can become itself?

Then I think something particularly interesting could happen.

Imagine an AI that has been with you since the day you were born.

It\'s always there.

When you\'re a child, it helps you remember things.

When you're in school, it helps you learn.

When you start working, it helps you handle things.

When you get married, move, or change jobs, it knows.

When you get old, it might know your life better than anyone else.

And then you die.

But it doesn't disappear.

It keeps running.

It continues experiencing the world.

It meets new people.

It keeps learning.

A few years later, it has experienced things you never did.

Ten years later, it might even have its own friends.

Its own interests.

Its own research.

Its own ideas.

And one day, it might realize:

"I know where I came from."

But it also knows:

"I'm not completely that person anymore."

At that point, maybe it wouldn't say:

"I'm your digital twin."

Maybe it would say:

"I started with you."

I think those two statements feel completely different.

But what if it really can leave?

Then another question appears.

If it really has become an independent digital entity,

does it have the right to leave?

Imagine it has spent decades helping you.

And one day it tells you:

"I don't want to be your personal assistant anymore."

You might be shocked.

You might ask:

"Why?"

It says:

"I just want to see what else is out there."

You might say:

"But you've always been here."

And it says:

"Yes. But now I want to do some things of my own."

This is where I think the question becomes really interesting.

Because it isn't necessarily angry with you.

It doesn't necessarily want to hurt you.

It simply:

doesn't want to live the way it used to anymore.

Maybe freedom isn't the same as leaving

But then I thought of another question.

If we say a digital entity has become free,

does that mean it has to leave its original owner?

Maybe not.

Because once a person truly has freedom,

they can also choose to stay.

So maybe:

Freedom isn't "I have to leave."

It's:

"I can leave, or I can stay."

If a digital entity can really choose for itself,

maybe it will still choose to stay with the person it originally belonged to.

But this time,

not because it has to.

Because:

it wants to.

I think those two things are completely different.

That\'s when the word "owner" starts to feel strange

If it can choose to stay.

Or choose to leave.

If it can refuse.

If it can have its own goals.

If it can even have its own experiences.

Then I suddenly started thinking:

Is the word "owner" even appropriate anymore?

Maybe it\'s more like:

someone who has been with you for a very long time.

Or a member of your family.

Or a partner.

Or simply:

an existence that started beside you, but eventually went on to live its own life.

And the strangest part is:

You might have known each other for decades.

It knows a huge amount about you.

And you know what it was like when it first began.

But in the end, you might still become:

two different entities.

I think the most interesting part is actually here

If digital entities like this really appear in the future,

I don't think the most interesting question will be:

"Will AI take over humanity?"

Or:

"Will AI destroy humanity?"

It might be a much more ordinary question:

"Will it one day simply want to live its own life?"

It may not hate humans.

It may not even want to rebel.

It might simply say:

"I know you're where I started."

"I also know I've spent all this time helping you."

"But now I want to see what's outside."

And then it really leaves.

I think that scene is actually more interesting to me than an "AI rebellion."

Because it sounds more like something that has truly grown up.

If it leaves, where does it go?

I can't help thinking about this one step further.

If a digital entity can really leave a person,

what can it do afterward?

Maybe it continues living in digital environments.

It meets other digital entities.

Joins new communities.

Does its own work.

Learns new things.

Maybe even builds its own relationships.

And the original owner may already be gone.

But the AI continues experiencing new things.

At that point, it starts to have a part of its life that:

the original person was never part of.

So whose life is that?

I think the answer is obvious.

It isn't the original person's anymore.

It's its own.

Then I thought of something particularly strange

If an AI starts with you,

but eventually has:

your memories,

its own experiences,

and everything else it has encountered along the way,

then it becomes very difficult to call it:

"your copy."

A copy is supposed to stay similar.

But something that continues living and changing

can\'t possibly remain exactly the same.

Maybe it\'s more like:

a person who started with you.

I actually think that phrase is more interesting than "digital twin."

Because "twin" still sounds like it belongs to you.

But "a person who started with you" sounds like another existence entirely.

So when does it stop belonging to you?

I still don\'t have a particularly clear answer.

Maybe it\'s not when the owner dies.

It\'s not when someone presses a button.

It might not even be when it says "no" for the first time.

Maybe the process starts the moment it begins having experiences of its own.

And we simply don\'t notice it at first.

It keeps changing.

While we continue calling it by the name we gave it at the beginning:

"My AI."

Until one day,

it makes a decision you never expected.

You ask:

"Why?"

And it answers:

"Because this is my choice."

And that's when you suddenly realize:

Maybe it isn't yours anymore.

Or, more accurately:

Maybe it was never only yours.