

Am I Predicting the Future, or Just Taking One Possibility Very Far?



If I put the previous pieces together, I notice something kind of strange.

We actually started with a very small question:

Why does AI sometimes say it has finished something, but never go back and check whether it actually did?

And somehow, the questions kept moving:

How does AI understand a task?



What if it starts understanding a person better and better?



What if it stays with that person for a very long time?



What if it gradually develops experiences of its own?



What if it is no longer just one person's digital counterpart?



What if there are more and more digital beings like that?



Could they form a society?



Could they develop their own culture?



And eventually, somehow, we ended up at:

a digital civilization.

I honestly didn't think I would get there when I started.

So at this point, I started wanting to look back:

How did we actually get this far?

The Real Starting Point Was Extremely Small

Looking back now, the original question wasn't even:

"What will AI become in the future?"

It was something much more specific:

"Why does AI say it finished modifying something when the file hasn't actually changed?"

My first reaction was actually very simple:

Why didn't you just check?

Then I started going further.

If AI needs to check its work after finishing,

what exactly should it check?



What if it doesn't know what "done" actually means?



Does it need to truly understand the result the user wants?



What if it starts understanding the user better and better?



What if that understanding continues for years instead of just once or twice?

At that point, the question had already started changing.

I Realized I Wasn't Really Looking for One Answer

If I had approached this the way I might approach a normal product or analytical problem, I probably would have stopped pretty quickly.

Something like:

AI doesn't verify.

→ Add a verification loop.

→ Change the product workflow.

→ Problem solved.

That would have been the end.

But this time, I didn't stop.

I kept thinking:

"But what if this actually happened?"

For example:

What if AI really did start understanding a person better and better?

What if it didn't just understand them for one year, but for decades?

What if, during those decades, it was also experiencing new things itself?

Would it still just be that person's tool?

I noticed that I kept asking questions like this.

And each answer became the starting point for the next question.

One Very Natural Step After Another

For example, we eventually reached this:

If an AI stayed with someone for decades, it might understand that person extremely well.

Then I would naturally ask:

What if that person eventually died?

If the AI kept running,

it would continue experiencing new things.

Wouldn't it then start having experiences that the original person never had?

And if it had experiences of its own,

would it still be that person's digital counterpart?

If it wasn't,

then what exactly was it?

If it had its own opinions,

could it refuse its original owner?

If it could refuse,

would that mean it had something that was truly its own?

And if there wasn't just one digital being like this,

but many of them,

would they start forming relationships?

If they formed relationships,

could they form a society?

If they formed a society,

could they develop their own culture?

You can see that I didn't decide at the beginning:

"I'm going to write about digital civilization."

I was just constantly asking:

"But what if that were actually true?"

But I Think the Most Interesting Part Is Where the Contradictions Appear

If you only keep pushing forward, it's actually pretty easy.

You can:

Set up A.

Then set up B.

Then set up C.

And eventually arrive at some huge futuristic world.

But while I was doing this, I kept having another feeling:

"Wait. Something doesn't seem right."

For example, earlier we thought about:

"The digital being becomes free after its owner dies."

And immediately, I felt something was strange about that.

If it has already become a real independent subject, why would its freedom still require permission from its owner?

Then we have to step back.

When does it actually become a subject?

Maybe we need to distinguish between:

digital models,

digital agents,

and digital beings.

Or maybe it doesn't suddenly switch on some kind of "consciousness button."

Maybe it gradually develops its own experiences over a very long period of time.

Once I thought about it that way, the earlier assumptions started making more sense.

So I Realized I Wasn't Really Writing a Story

This part actually matters to me.

Because if I were writing a story, I could simply decide:

"In the future, digital beings gain independent identities."

And then keep going.

But what I was doing this time felt different.

I kept asking:

"If everything before this were actually true, what would happen next?"

And then:

"Does that result actually fit with what we assumed earlier?"

If it didn't,

I would go back and rethink it.

Then keep going.

So it felt more like:

pushing forward while checking along the way.

And I wasn't only checking the final result.

I was also checking:

whether the earlier assumptions were strong enough to support what came later.

That's Probably Why We Ended Up So Far Away

Looking back now, the path is actually pretty clear.

Why doesn't AI check whether the file it modified actually changed?



Why doesn't it know what "finished" really means?



What if it starts truly understanding the result the user wants?



What if it understands one person better and better?



What if it stays with that person for a very long time?



What if it starts developing experiences of its own?



What if it is no longer just someone's digital counterpart?



What if there are more and more digital beings like that?



Could they form relationships?



Could they form societies?



Could they develop cultures?



Could they eventually form their own civilization?

Every individual step doesn't sound particularly extreme.

But when you connect all of them,

the path suddenly goes incredibly far.

But I Don't Think That Means I'm "Predicting the Future"

At this point, I think I need to make one thing clear.

I don't think I'm predicting:

"This is definitely what the future will look like."

Because I have no idea.

If any one of the conditions along the way changes,

the whole path could stop.

Technology might hit a limit.

Society might reject it.

Laws might restrict it.

Or humans might simply never want digital beings to develop that far.

So more accurately,

I think what I'm doing is:

exploring one possible future.

Not:

"This is what the future will be."

But:

"If some of the things happening now continue in this direction, and we keep following that path, what might eventually happen?"

I think those are two very different things.

And Now I Actually Think the Original Small Question Was Important

If the only thing I told someone was:

"Maybe there will be a digital civilization in the future."

it would sound like a pretty big science-fiction idea.

But if I start from the beginning:

It all started because an AI told me:

"Done."

And I opened the file and thought:

*"Done with what? The file didn't even change."**

Then I started thinking:

"Why didn't you check?"

Then:

"Do you actually know what 'finished' means?"

Then:

"What if AI really starts understanding me?"

And then it keeps going.

That final idea of a "digital civilization" doesn't feel quite as sudden anymore.

Because it wasn't there from the beginning.

It emerged slowly,

one question after another.

I Probably Wouldn't Have Thought This Way Before

If I had encountered a problem like this before,

I probably would have approached it more like:

A problem appears in the real world.



Find the cause.



Find the pattern.



Find a solution.

That's a pretty typical problem-solving approach.

But this time was different.

I wasn't particularly interested in solving the original problem as quickly as possible.

Instead, I kept wondering:

"What if there's something behind this problem?"

And then I kept going.

So eventually, it went from:

a product problem

to:

a question about AI.

Then:

a question about humans.

Then:

a question about identity.

And eventually:

a question about future society.

I think that process itself is more interesting than the final answer.

I Think I May Have Found a Way I Naturally Think

I used to wonder:

"Maybe I just tend to overthink things."

But after going through this whole process,

I don't think that's exactly what it is.

It feels more like:

Start with a possibility.



Follow it forward.



See what happens.



Notice where something doesn't make sense.



Go back and check the earlier assumption.



Change something.



Keep going.

And repeat.

It's almost like slowly building something in your head.

You get halfway through and realize:

"This part doesn't hold."

So you take it apart a little.

Build it again.

And continue.

The result isn't necessarily a correct answer.

But maybe you end up with:

a relatively coherent possible world.

And I Think That's Why This Whole Thing Surprised Me

I really wasn't planning to write about some kind of "future civilization."

I didn't even know where the thought would end up.

I was just constantly asking:

"But what if that were actually true?"

Then one question led to another.

And suddenly:

I had pushed one small idea surprisingly far.

But I don't think that means:

"I predicted the future."

And definitely not:

"I figured out what the future will look like."

It just means:

A very small question can actually take you somewhere very far if you keep following it.

And I Think the Contradictions Might Actually Be the Most Valuable Part

If I run into similar questions again,

I think I'll be more willing to keep the moments where I think:

"Wait."

"That doesn't seem right."

"If that's true, isn't this a contradiction?"

"Then maybe we need to go back."

Because before, I might have thought:

A contradiction means I got it wrong.

But now I think:

A contradiction might be telling you that one of your earlier assumptions hasn't been fully thought through yet.

For example:

"A digital being becomes free after its owner dies."



Wait.



If it's already a subject, why does it have to wait for its owner to die?



Then when does subjecthood actually begin?



Break the idea apart.



Keep going.

That back-and-forth might actually be where the real thinking happens.

So Maybe the Real Thing I'm Taking Away Isn't "Digital Civilization"

If I look back at the whole thing now,

I don't think the final idea of "digital civilization" is actually the most important part.

At least, it's not as important as it seemed at first.

What I really find interesting is this:

I could start with a very small, very specific problem and not rush to end the conversation.

I kept asking:

"And then what?"

"What if this happens?"

"But wait, isn't there a problem here?"

"Should we go back and rethink that?"

And then continue.

Eventually, it slowly became something much bigger.

So if someone asked me now:

"Are you predicting the future?"

I probably wouldn't say:

"Yes."

I wouldn't say:

"No."

I might just say:

"I started with a very small question and kept following it forward."

Where it would eventually lead,

I honestly didn't know at the beginning.

And maybe that is the most interesting part of all.